

CGM data analysis report

Xinyu Dong

2026-01-18

Download data

The *CGMdata.csv* dataset is a perturbed and truncated toy dataset of Continuous Glucose Measurements based (rather loosely) off real-world readings of glucose levels of patients residing in the Emory healthcare system. In real world scenarios, though the patient is monitored closely during their occupancy in a hospital, due to technical reasons the glucose readings may be partially *missing*. Ideally we could predict the start of these missing intervals and perform data imputation as a next step.

Variables in the original dataset include:

- patient_id: the encoded patient id of a patient, patient id differs across patients and are categorical.
- glucose: a single numeric continuous glucose reading of a patient at time t , where t is indicated by the variable *time*.
- time: the time of the glucose reading in minutes from baseline observation.
- x1: an unknown feature variable in numerical form. Its value remains constant in patient, and could differ across patients.
- x2: an unknown feature variable in numerical form. Its value remains constant in patient, and could differ across patients.

For simplicity, this analysis focuses on variables patient_id, glucose, and time only. The research question of this analysis is: Can we use a reasonable statistical method to accurately predict the start of a missing interval in the glucose readings of patients? What is the accuracy of prediction?

Data Download:

```
# first we load the required packages
library(here)
```

```
## here() starts at C:/Users/xdong63/OneDrive - Emory/Desktop/Fall 2024 Emory/BIOS 731/BIOS 731 homework
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```

## The following objects are masked from 'package:stats':
##
##     filter, lag

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

library(scales)

## Warning: package 'scales' was built under R version 4.4.3

library(tidyr)
library(ggplot2)

## Warning: package 'ggplot2' was built under R version 4.4.3

# then we read in the data using a "here" (relative path statement)
CGM_data <- read.csv(here("data", "CGMdata.csv"))
CGM_data <- CGM_data[, 1:3]

```

Data Cleaning

Now that we downloaded the continuous glucose monitoring data and filtered out the variables of interest, we need to fill in the missing time stamps first: glucose level for all missing timestamps are set to be zero as an indication of unobserved data, and we create a new indicator variable *ind* to indicate whether the data is observed (value 1) or missing (value 0).

Another indicator variable *start_ind* serves as the parameter of interest, which indicate whether the data is the start of a missing interval (value 1) or not (value 0). The time gap between two lines of data within the same individual is enforces to be 5 (minutes), which serves as the criteria for whether we should “insert” a line of data between two observed readings or not.

```

if(!file.exists((here::here("data", "CGMdata_clean.csv")))){
  source(here("data", "data_cleaning.R"))
} else{
  readr::read_csv(here::here("data", "CGMdata_clean.csv"),
  show_col_types = FALSE)
}

## # A tibble: 26,365 x 5
##   patient_id  glucose    time    ind start_ind
##   <chr>       <dbl> <dbl> <dbl>    <dbl>
## 1 EM008_extra     253     0     1        0
## 2 EM008_extra     253     5     1        0
## 3 EM008_extra     257    10     1        0
## 4 EM008_extra     263    15     1        0
## 5 EM008_extra     244    20     1        0
## 6 EM008_extra     225    25     1        0
## 7 EM008_extra     224    30     1        0
## 8 EM008_extra     231    35     1        0
## 9 EM008_extra     235    40     1        0
## 10 EM008_extra    243    45     1        0
## # i 26,355 more rows

```

Data Modeling

Let Y_i denote the binary outcome for observation i , where

$$Y_i = \begin{cases} 1, & \text{if a start event occurs (\texttt{start_ind} = 1)} \\ 0, & \text{otherwise (\texttt{start_ind} = 0).} \end{cases}$$

Let $\mathbf{x}_i = (1, x_{i1}, x_{i2}, \dots, x_{ip})^\top$ be the vector of predictors for observation i , including an intercept term. In this analysis, all available feature variables are considered as predictors except `ind`, `x1`, `x2`, which are excluded from modeling.

Model specification

We model the conditional probability of a start event using a binary logistic regression:

$$\pi_i \equiv \Pr(Y_i = 1 \mid \mathbf{x}_i),$$

with the logit link function

$$\log\left(\frac{\pi_i}{1 - \pi_i}\right) = \mathbf{x}_i^\top \boldsymbol{\beta} = \beta_0 + \sum_{j=1}^p \beta_j x_{ij}.$$

Equivalently, the predicted probability is

$$\pi_i = \frac{\exp(\mathbf{x}_i^\top \boldsymbol{\beta})}{1 + \exp(\mathbf{x}_i^\top \boldsymbol{\beta})}.$$

Estimation

Parameters $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^\top$ are estimated by maximum likelihood. Under independence across observations, the likelihood is

$$L(\boldsymbol{\beta}) = \prod_{i=1}^n \pi_i^{Y_i} (1 - \pi_i)^{1 - Y_i},$$

and the log-likelihood is

$$\ell(\boldsymbol{\beta}) = \sum_{i=1}^n [Y_i \log(\pi_i) + (1 - Y_i) \log(1 - \pi_i)].$$

The maximum likelihood estimate $\hat{\boldsymbol{\beta}}$ is obtained by numerically maximizing $\ell(\boldsymbol{\beta})$.

Train/validation evaluation

To assess generalization performance, we split the dataset into a training set and a validation set. Model parameters are estimated using only the training data, and predictions and performance metrics are computed on the held-out validation data. This procedure reduces optimistic bias compared to evaluating on the same data used for model fitting.

Prediction and classification

For each observation i , we compute the fitted probability

$$\hat{\pi}_i = \frac{\exp(\mathbf{x}_i^\top \hat{\boldsymbol{\beta}})}{1 + \exp(\mathbf{x}_i^\top \hat{\boldsymbol{\beta}})}.$$

To obtain a hard class label, we apply a decision threshold $\tau = 0.5 \in (0, 1)$:

$$\widehat{Y}_i = \begin{cases} 1, & \text{if } \widehat{\pi}_i \geq \tau, \\ 0, & \text{if } \widehat{\pi}_i < \tau. \end{cases}$$

Using the predicted labels \widehat{Y}_i and true labels Y_i , we summarize performance with a confusion matrix and compute accuracy:

$$\text{Accuracy} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(\widehat{Y}_i = Y_i),$$

where $\mathbb{I}(\cdot)$ is the indicator function.

```
if(!file.exists(here::here("data", "CGMdata_pred.csv"))){
  source(here("modeling", "CGM_data_modeling.R"))
} else{
  readr::read_csv(here::here("data", "CGMdata_pred.csv"),
    show_col_types = FALSE)
}

## # A tibble: 7,910 x 7
##   patient_id  glucose  time  ind start_ind pred_prob pred_class
##   <chr>        <dbl>  <dbl> <dbl>     <dbl>      <dbl>       <dbl>
## 1 EM008_extra     253     0     1        0  2.22e-16       0
## 2 EM008_extra     257    10     1        0  2.22e-16       0
## 3 EM008_extra     263    15     1        0  2.22e-16       0
## 4 EM008_extra     244    20     1        0  2.22e-16       0
## 5 EM008_extra     235    40     1        0  2.22e-16       0
## 6 EM008_extra     243    45     1        0  2.22e-16       0
## 7 EM008_extra     233    65     1        0  2.22e-16       0
## 8 EM008_extra     220   100     1        0  2.22e-16       0
## 9 EM008_extra     220   105     1        0  2.22e-16       0
## 10 EM008_extra    215   140     1        0  2.22e-16      0
## # i 7,900 more rows

CGM_data_pred <- readr::read_csv(here::here("data", "CGMdata_pred.csv"),
  show_col_types = FALSE)

if (!exists("accuracy_val") || !exists("cm")) {
  accuracy_val <- mean(CGM_data_pred$pred_class == CGM_data_pred$start_ind, na.rm = TRUE)
  cm <- table(Truth = CGM_data_pred$start_ind, Prediction = CGM_data_pred$pred_class)
}

cat("\nValidation accuracy:", round(accuracy_val, 4), "\n\n")

##
## Validation accuracy: 0.9987

print(cm)

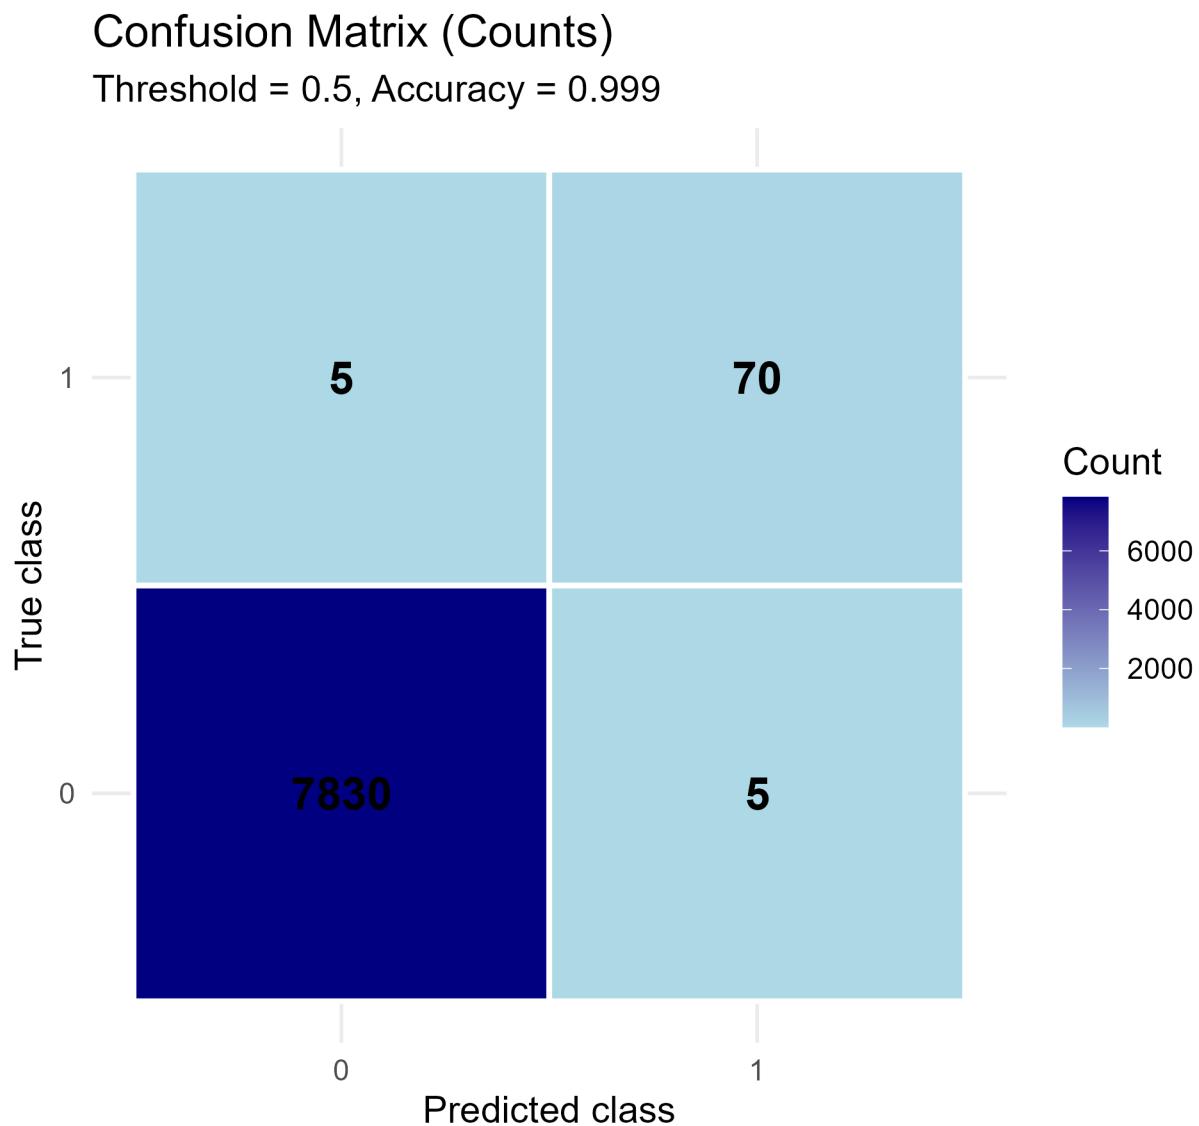
##          Prediction
## Truth      0      1
##      0 7830      5
##      1      5     70
```

Visualization

We perform a visualization of the confusion matrix obtained in section **Modeling**. The first plot shows a confusion matrix by raw counts, and the second shows a row normalized confusion matrix.

```
source(here("visualization", "CGM_data_visualization.R"))
```

```
knitr::include_graphics(  
  here("results", "confusion_matrix_counts.png")  
)
```



```
knitr::include_graphics(  
  here("results", "confusion_matrix_row_normalized.png")  
)
```

Confusion Matrix (Row-normalized)

Threshold = 0.5, Accuracy = 0.999

